

BHM_example

February 17, 2022

NRT Lectures - Statistical Modeling

1 Bayesian Hierarchical Model

1.0.1 Rat Tumor Example

```
[1]: import math
import random
import numpy as np
import pandas as pd
# import graphviz
# from pymc3 import model_to_graphviz
import pymc3 as pm
from pymc3 import Model, sample, Beta, Binomial, Exponential, Uniform, summary,
    plot_posterior, model_to_graphviz, Deterministic
import matplotlib.pyplot as plt
# import os
# os.environ["PATH"] += os.pathsep + 'C:\Program
    Files\Python37\Lib\site-packages\graphviz\dot.py'
```

```
[2]: d = pd.read_table("rattumor.txt", sep = " ")
d = d.iloc[:, :2]
d
```

```
[2]:
```

	y	N
0	0	20
1	0	20
2	0	20
3	0	20
4	0	20
..
66	16	52
67	15	46
68	15	47
69	9	24
70	4	14

[71 rows x 2 columns]

```
[3]: d.describe()
```

```
[3]:
```

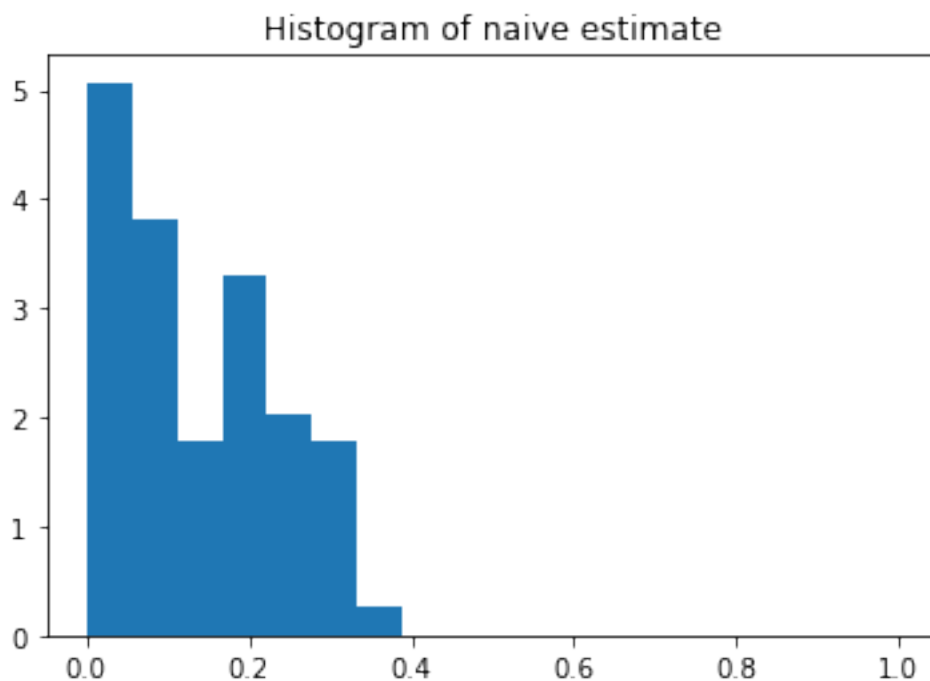
	y	N
count	71.000000	71.000000
mean	3.760563	24.492958
std	3.811504	10.973830
min	0.000000	10.000000
25%	1.000000	19.000000
50%	3.000000	20.000000
75%	5.000000	22.500000
max	16.000000	52.000000

A naive estimate of θ_j is $\hat{\theta}_j = y_j/n_j$

Histogram of $\hat{\theta}_j$

```
[4]: plt.hist(d.y/d.N, range = (0,1), bins = 18, density=True)
plt.title("Histogram of naive estimate")
plt.show
```

```
[4]: <function matplotlib.pyplot.show(close=None, block=None)>
```



```
[5]: with Model() as model1:

    # Priors
```

```

alpha = Exponential('alpha', 0.001)
beta = Exponential('beta', 0.001)

theta = Beta('theta', alpha=alpha, beta=beta, shape=71)

# Data likelihood
y_like = Binomial('y_like', n=d.N, p=theta, observed=d.y)

```

```

[6]: random.seed(100)
      with model1:
          trace1 = sample(100, tune=100)

```

/tmp/ipykernel_9952/1044724547.py:3: FutureWarning: In v4.0, pm.sample will return an `arviz.InferenceData` object instead of a `MultiTrace` by default. You can pass return_inferencedata=True or return_inferencedata=False to be safe and silence this warning.

```

      trace1 = sample(100, tune=100)
Only 100 samples in chain.
Auto-assigning NUTS sampler...
Initializing NUTS using jitter+adapt_diag...
Multiprocess sampling (2 chains in 2 jobs)
NUTS: [theta, beta, alpha]

```

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

Sampling 2 chains for 100 tune and 100 draw iterations (200 + 200 draws total) took 3 seconds.

The acceptance probability does not match the target. It is 0.9763936443680605, but should be close to 0.8. Try to increase the number of tuning steps.

The acceptance probability does not match the target. It is 0.920349624272922, but should be close to 0.8. Try to increase the number of tuning steps.

The rhat statistic is larger than 1.05 for some parameters. This indicates slight problems during sampling.

The number of effective samples is smaller than 10% for some parameters.

```

[7]: summary(trace1)

```

/home/ooblack/miniconda3/envs/viz/lib/python3.10/site-packages/arviz/data/io_pymc3.py:96: FutureWarning: Using `from_pymc3` without the model will be deprecated in a future release. Not using the model will return less accurate and less useful results. Make sure you use the model argument or call from_pymc3 within a model context.

```

      warnings.warn(

```

```

[7]:
      mean      sd  hdi_3%  hdi_97%  mcse_mean  mcse_sd  ess_bulk  \
alpha      3.119  1.399   1.164   5.611    0.368    0.272    17.0
beta     18.891  7.971   6.632  34.316    1.771    1.271    22.0
theta[0]   0.075  0.045   0.001   0.163    0.005    0.004    58.0

```


theta[1]	0.071	0.038	0.003	0.135	0.005	0.003	56.0
theta[2]	0.069	0.045	0.004	0.150	0.004	0.003	119.0
...
theta[66]	0.259	0.056	0.160	0.357	0.003	0.002	349.0
theta[67]	0.274	0.058	0.170	0.384	0.003	0.003	419.0
theta[68]	0.265	0.057	0.151	0.371	0.003	0.003	304.0
theta[69]	0.267	0.079	0.151	0.444	0.005	0.004	288.0
theta[70]	0.201	0.066	0.094	0.324	0.003	0.002	285.0

	ess_tail	r_hat
alpha	21.0	1.09
beta	33.0	1.07
theta[0]	63.0	1.03
theta[1]	73.0	1.03
theta[2]	87.0	1.01
...
theta[66]	121.0	1.01
theta[67]	157.0	1.05
theta[68]	141.0	1.02
theta[69]	118.0	1.08
theta[70]	198.0	1.00

[73 rows x 9 columns]


```
[8]: plot_posterior(trace1['alpha'])
```

```
[8]: <AxesSubplot:title={'center':'x'}>
```




```
[9]: plot_posterior(trace1['beta'])
```


```
[9]: <AxesSubplot:title={'center':'x'}>
```



```
[10]: alpha = trace1.get_values(varname='alpha')
      beta = trace1.get_values(varname='beta')
      plt.scatter(alpha, beta)
      plt.xlabel('alpha')
      plt.ylabel('beta')
      plt.show()
```



```
[11]: plt.scatter(np.log(alpha/beta), np.log(alpha+beta))
plt.xlabel('log(alpha/beta)')
plt.ylabel('log(alpha+beta)')
plt.show()
```



1.0.2 Try another prior

```
[12]: with Model() as model2:
```

```
    phi1 = Uniform('phi1', lower=0, upper=1)
    phi2 = Uniform('phi2', lower=0, upper=1000)

    alpha = Deterministic('alpha', phi1 / (phi2**2))
    beta = Deterministic('beta', (1-phi1) / phi2**2)

    theta = Beta('theta', alpha=alpha, beta=beta, shape=71)

    # Data likelihood
    y_like = Binomial('y_like', n=d.N, p=theta, observed=d.y)
```

```
[13]: random.seed(100)
      with model2:
          trace2 = sample(100, tune=100)
```

/tmp/ipykernel_9952/2992495953.py:3: FutureWarning: In v4.0, pm.sample will return an `arviz.InferenceData` object instead of a `MultiTrace` by default. You can pass return_inferencedata=True or return_inferencedata=False to be safe and silence this warning.

```
    trace2 = sample(100, tune=100)
```

Only 100 samples in chain.

Auto-assigning NUTS sampler...

Initializing NUTS using jitter+adapt_diag...

Multiprocess sampling (2 chains in 2 jobs)

NUTS: [theta, phi2, phi1]

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

Sampling 2 chains for 100 tune and 100 draw iterations (200 + 200 draws total) took 5 seconds.

The acceptance probability does not match the target. It is 0.9346941703130048, but should be close to 0.8. Try to increase the number of tuning steps.

The acceptance probability does not match the target. It is 0.9318999435860623, but should be close to 0.8. Try to increase the number of tuning steps.

The rhat statistic is larger than 1.05 for some parameters. This indicates slight problems during sampling.

The number of effective samples is smaller than 25% for some parameters.

```
[14]: summary(trace2)
```

/home/ooblack/miniconda3/envs/viz/lib/python3.10/site-

packages/arviz/data/io_pymc3.py:96: FutureWarning: Using `from_pymc3` without the model will be deprecated in a future release. Not using the model will return less accurate and less useful results. Make sure you use the model argument or call from_pymc3 within a model context.

warnings.warn(

```
[14]:
```


	mean	sd	hdi_3%	hdi_97%	mcse_mean	mcse_sd	ess_bulk	\
phi1	0.144	0.013	0.119	0.167	0.001	0.001	182.0	
phi2	0.252	0.038	0.184	0.320	0.006	0.004	43.0	
alpha	2.414	0.730	1.190	3.703	0.120	0.086	35.0	
beta	14.438	4.277	6.428	21.429	0.605	0.430	50.0	
theta[0]	0.061	0.041	0.006	0.148	0.003	0.002	159.0	
...		
theta[66]	0.266	0.059	0.166	0.363	0.003	0.002	442.0	
theta[67]	0.278	0.055	0.171	0.381	0.003	0.002	460.0	
theta[68]	0.275	0.061	0.175	0.381	0.003	0.002	384.0	
theta[69]	0.282	0.065	0.179	0.411	0.003	0.002	371.0	
theta[70]	0.215	0.074	0.081	0.345	0.004	0.003	296.0	

	ess_tail	r_hat
phi1	119.0	1.00
phi2	217.0	1.04
alpha	149.0	1.05
beta	217.0	1.04
theta[0]	132.0	1.00
...
theta[66]	138.0	1.00
theta[67]	162.0	1.02
theta[68]	230.0	0.99
theta[69]	124.0	1.00
theta[70]	181.0	1.00

[75 rows x 9 columns]


```
[15]: plot_posterior(trace2['alpha'])
```

```
[15]: <AxesSubplot:title={'center':'x'}>
```



```
[16]: plot_posterior(trace2['beta'])
```

```
[16]: <AxesSubplot:title={'center': 'x'}>
```



```
[17]: alpha2 = trace2.get_values(varname='alpha')
      beta2 = trace2.get_values(varname='beta')
      plt.scatter(np.log(alpha2/beta2), np.log(alpha2+beta2))
      plt.xlabel('log(alpha/beta)')
      plt.ylabel('log(alpha+beta)')
      plt.show()
```

